


```
x_train, x_test, y_train, y_test = train_test_split(x_train, y_train, test_size=0.2, stratify=y_train)
x_train_cv, x_test_cv, y_train_cv, y_test_cv = train_test_split(x_train, y_train, test_size=0.2, stratify=y_train)
```

```
In [45]: x_train.shape
Out[45]: (3267, 10)
In [46]: x_test.shape
Out[46]: (1021, 10)
In [47]: y_train.value_counts()
Out[47]:
0    3108
1     153
Name: stroke, dtype: int64
In [48]: rom = RandomOverSampler(random_state=42)
x_train_cv, y_train_cv = rom.fit_resample(x_train, y_train)
In [49]: x_train.shape
Out[49]: (6216, 10)
In [50]: y_train.value_counts()
Out[50]:
0    3108
1     153
Name: stroke, dtype: int64
```

One-Hot Encoding of Categorical Data

```
In [51]: ohe = ColumnTransformer([('ohe', OneHotEncoder(handle_unknown='ignore'), ['gender', 'ever_married', 'work_type', 'train_type', 'train_shape', y_train.shape])], remainder='passthrough')
x_train_ohe = ohe.transform(x_train)
x_cv_ohe = ohe.transform(x_cv)
x_test_ohe = ohe.transform(x_test)
print('After Vectorization.....')
print(x_train_ohe.shape, y_train.shape)
print(x_cv_ohe.shape, y_cv.shape)
print(x_test_ohe.shape, y_test.shape)
(6216, 16) (6216, 1)
After Vectorization.....
(6216, 21) (6216, 1)
(817, 21) (817, 1)
(1021, 21) (1021, 1)
In [52]: x_train_ohe = x_train_ohe[:, :16]
x_cv_ohe = x_cv_ohe[:, :16]
x_test_ohe = x_test_ohe[:, :16]
print('After Vectorization.....')
print(x_train_ohe.shape, y_train.shape)
print(x_cv_ohe.shape, y_cv.shape)
print(x_test_ohe.shape, y_test.shape)
(6216, 16) (6216, 1)
(817, 16) (817, 1)
(1021, 16) (1021, 1)
In [53]: features = ohe.get_feature_names()
features
Out[53]:
'ohe_x0_Female',
'ohe_x0_Male',
'ohe_x0_Other',
'ohe_x1_No',
'ohe_x1_Yes',
'ohe_x2_Govt_Job',
'ohe_x2_Never_Worked',
'ohe_x2_Private',
'ohe_x2_Self-employed',
'ohe_x3_Childrent',
'ohe_x3_Rural',
'ohe_x3_Urban',
'ohe_x4_Unknown',
'ohe_x4_formerly smoked',
'ohe_x4_Never smoked',
'ohe_x4_smokes',
'age',
'hypertension',
'heart_disease',
'avag_glucose_level',
'tmi'
```

```
In [54]: x_train_hyp = np.array(x_train['hypertension']).reshape((-1,1))
x_cv_hyp = np.array(x_cv['hypertension']).reshape((-1,1))
x_test_hyp = np.array(x_test['hypertension']).reshape((-1,1))
print('After Vectorization.....')
print(x_train_hyp.shape, y_train.shape)
print(x_cv_hyp.shape, y_cv.shape)
print(x_test_hyp.shape, y_test.shape)
(6216, 1) (6216, 1)
(817, 1) (817, 1)
(1021, 1) (1021, 1)
In [55]: x_train_hd = np.array(x_train['heart_disease']).reshape((-1,1))
x_cv_hd = np.array(x_cv['heart_disease']).reshape((-1,1))
x_test_hd = np.array(x_test['heart_disease']).reshape((-1,1))
print('After Vectorization.....')
print(x_train_hd.shape, y_train.shape)
print(x_cv_hd.shape, y_cv.shape)
print(x_test_hd.shape, y_test.shape)
(6216, 1) (6216, 1)
(817, 1) (817, 1)
(1021, 1) (1021, 1)
In [56]: std = ColumnTransformer([('norm', MinMaxScaler(), ['age', 'avag_glucose_level', 'tmi'])], remainder='drop')
std.fit(x_train)
x_train_std = std.transform(x_train)
x_cv_std = std.transform(x_cv)
x_test_std = std.transform(x_test)
print('After Vectorization.....')
print(x_train_std.shape, y_train.shape)
print(x_cv_std.shape, y_cv.shape)
print(x_test_std.shape, y_test.shape)
(6216, 10) (6216, 1)
(6216, 3) (6216, 1)
(817, 3) (817, 1)
(1021, 3) (1021, 1)
```

Combining all encoded columns

```
In [57]: x_tr = np.hstack((x_train_ohe.astype(np.float), x_train_hyp.astype(np.float), x_train_hd.astype(np.float), x_train_std.astype(np.float), x_train_cv_ohe.astype(np.float), x_cv_hyp.astype(np.float), x_cv_hd.astype(np.float), x_cv_std.astype(np.float), x_test_ohe.astype(np.float), x_test_hyp.astype(np.float), x_test_hd.astype(np.float), x_test_std.astype(np.float)))
print('Final Data Matrix Shape is.....')
print(x_tr.shape, y_train.shape)
print(x_cv.shape, y_cv.shape)
print(x_te.shape, y_test.shape)
Final Data Matrix Shape is.....
(6216, 21) (6216, 1)
(817, 21) (817, 1)
(1021, 21) (1021, 1)
In [58]: def cmf_matrix(true_y, pred_y):
cmf_matrix = confusion_matrix(y_test, predicted_y)
print('cmf_matrix', cmf_matrix)
group_percentages = [(0.25, 0.75).format(value) for value in cmf_matrix.flatten()]
labels = ['t', 'p']
sns.heatmap(cmf_matrix, annot=True, cmap='YlGnBu', fmt='.3f', xticklabels=[0,1], yticklabels=[0,1])
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
# Precision Matrix
pc_matrix = cmf_matrix/cmf_matrix.sum(axis=0)
print('pc_matrix', pc_matrix)
sns.heatmap(pc_matrix, annot=True, cmap='YlGnBu', fmt='.3f', xticklabels=[0,1], yticklabels=[0,1])
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
# Recall Matrix
rl_matrix = ((cmf_matrix.T)/(cmf_matrix.sum(axis=1))).T
print('rl_matrix', rl_matrix)
sns.heatmap(rl_matrix, annot=True, cmap='YlGnBu', fmt='.3f', xticklabels=[0,1], yticklabels=[0,1])
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```

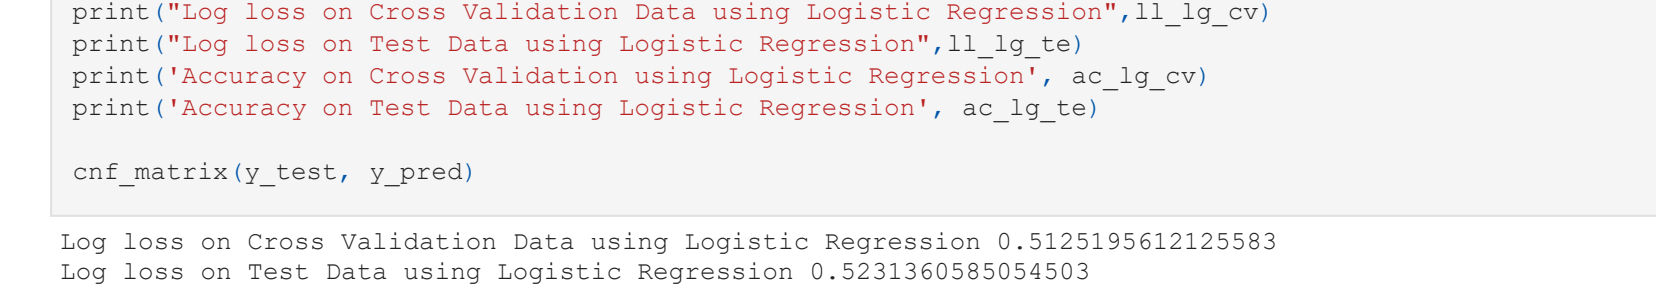
Random Model and its Performance

```
In [59]: test_data_len = x_test.shape[0]
cv_data_len = x_cv.shape[0]
# We create a output array that has exactly same size as the cv data
cv_predicted_y = np.zeros((cv_data_len,2))
for i in range(cv_data_len):
    rand_probs = np.random.rand(1,2)
    cv_predicted_y[i] = (rand_probs/sum(sum(rand_probs)))
# Test-Set Error:
# We create a output array that has exactly same as the test data
test_predicted_y = np.zeros((test_data_len,2))
for i in range(test_data_len):
    rand_probs = np.random.rand(1,2)
    test_predicted_y[i] = (rand_probs/sum(sum(rand_probs)))
predicted_y = np.argmax(test_predicted_y, axis=1)
predicted_cv = np.argmax(cv_predicted_y, axis=1)
ll_rm_cv = log_loss(y_cv, cv_predicted_y, eps=1e-15)
ac_rm_cv = accuracy_score(y_cv, predicted_cv)
ll_rm_te = log_loss(y_test, test_predicted_y, eps=1e-15)
ac_rm_te = accuracy_score(y_test, predicted_y)
print('Log loss on Cross Validation Data using Random Model', ll_rm_cv)
print('Log loss on Test Data using Random Model', ll_rm_te)
print('Accuracy on Cross Validation using Random Model', ac_rm_cv)
print('Accuracy on Test Data using Random Model', ac_rm_te)
cmf_matrix(y_test, predicted_y)
```


Log loss on Cross Validation Data using Random Model 0.9264253972091854
Log loss on Test Data using Random Model 0.93267590369544
Accuracy on Cross Validation using Random Model 0.749092007349413
Accuracy on Test Data using Random Model 0.4799216454564153



Precision matrix (Column Sum=1)



Recall matrix (Row sum=1)



Logistic Regression

Hyperparameter Tuning

```
In [60]: alpha = [10 ** x for x in range(-6, 4)]
params = {'alpha':alpha}
clf1 = SGDClassifier(loss='log', n_jobs=-1, random_state=42)
r_search = RandomizedSearchCV(clf1, param_distributions=params, return_train_score=True, random_state=42)
r_search.fit(x_tr, y_train)
Out[60]: RandomizedSearchCV(estimator=SGDClassifier(loss='log', n_jobs=-1, random_state=42),
param_distributions={'alpha': [1e-06, 1e-05, 0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000]},
random_state=42, return_train_score=True)
In [61]: print(f'The best hyperparameter values is {r_search.best_params_} at which the score is {r_search.best_score}')
The best hyperparameter values is {'alpha': 0.001} at which the score is 0.7689837365775675
```

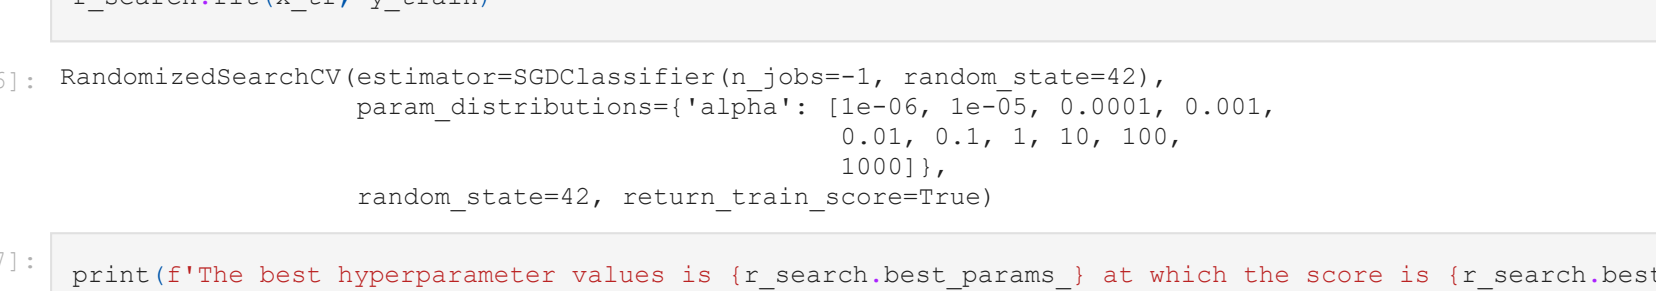
Training the model

```
In [62]: clf1 = SGDClassifier(loss='log', n_jobs=-1, random_state=42, **r_search.best_params_)
clf1.fit(x_tr, y_train)
cal_clf1 = CalibratedClassifierCV(clf1, cv='prefit')
cal_clf1.fit(x_tr, y_train)
Out[62]: CalibratedClassifierCV(base_estimator=SGDClassifier(alpha=0.001, loss='log', n_jobs=-1, random_state=42),
cv='prefit')
In [63]: y_pred_cv = cal_clf1.predict(x_cv)
y_prob_cv = cal_clf1.predict_proba(x_cv)
y_pred = cal_clf1.predict(x_te)
y_prob = cal_clf1.predict_proba(x_te)
```

Performance of the model

```
In [64]: ll_lg_cv = log_loss(y_cv, y_prob_cv, eps=1e-15)
ac_lg_cv = accuracy_score(y_cv, y_pred_cv)
ll_lg_te = log_loss(y_test, y_prob, eps=1e-15)
ac_lg_te = accuracy_score(y_test, y_pred)
print('Log loss on Cross Validation Data using Logistic Regression', ll_lg_cv)
print('Log loss on Test Data using Logistic Regression', ll_lg_te)
print('Accuracy on Cross Validation using Logistic Regression', ac_lg_cv)
print('Accuracy on Test Data using Logistic Regression', ac_lg_te)
cmf_matrix(y_test, y_pred)
```

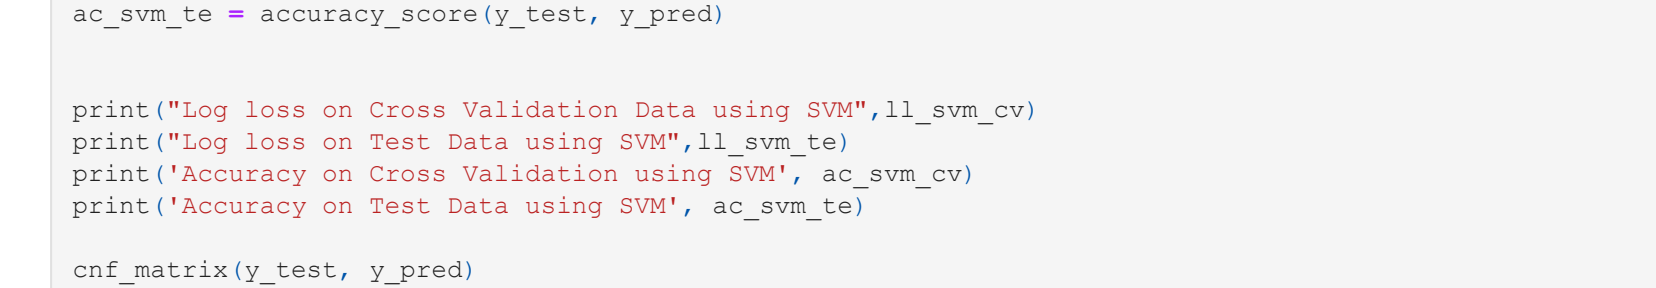
Log loss on Cross Validation Data using Logistic Regression 0.5125195612122583
Log loss on Test Data using Logistic Regression 0.5231360595054503
Accuracy on Cross Validation using Logistic Regression 0.749092007349413
Accuracy on Test Data using Logistic Regression 0.732615083251714



Precision matrix (Column Sum=1)




Recall matrix (Row sum=1)



Feature Importance

```
In [65]: importance = clf1.coef_
# summarize feature importance
for i, v in enumerate(importance[0]):
    print(f'Feature: {i}, Score: {v}')
plt.figure(figsize=(20,7))
sns.barplot(x=[x for x in range(importance.shape[1])], y=importance[0]).set_xticklabels(features, rotation=90)
plt.show()
Feature: 0, Score: 0.08744457495307384
Feature: 1, Score: 0.0811951641250333
Feature: 2, Score: -0.01814649413473182
Feature: 3, Score: -0.017645247504174698
Feature: 4, Score: 0.16861628935379094
Feature: 5, Score: -0.264252568149693
Feature: 6, Score: -0.0806056319606256
Feature: 7, Score: 0.227103133535398
Feature: 8, Score: -0.2002913243679036
Feature: 9, Score: 0.4684623711092656
Feature: 10, Score: 0.061762321019317
Feature: 11, Score: 0.0862398982100894
Feature: 12, Score: 0.02705792566014
Feature: 13, Score: 0.1744792935278436
Feature: 14, Score: -0.0939924720277727
Feature: 15, Score: 0.091033642147913
Feature: 16, Score: 0.4312492605722929
Feature: 17, Score: 0.346302157882967
Feature: 18, Score: 5.478478846717
Feature: 19, Score: 0.595938424235832
Feature: 20, Score: -0.05684716975492406
```



Support Vector Machines

Hyperparameter Tuning

```
In [66]: alpha = [10 ** x for x in range(-6, 4)]
params = {'alpha':alpha}
clf2 = SGDClassifier(loss='hinge', n_jobs=-1, random_state=42)
r_search = RandomizedSearchCV(clf2, param_distributions=params, return_train_score=True, random_state=42)
r_search.fit(x_tr, y_train)
Out[66]: RandomizedSearchCV(estimator=SGDClassifier(n_jobs=-1, random_state=42),
param_distributions={'alpha': [1e-06, 1e-05, 0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000]},
random_state=42, return_train_score=True)
In [67]: print(f'The best hyperparameter values is {r_search.best_params_} at which the score is {r_search.best_score}')
The best hyperparameter values is {'alpha': 0.0001} at which the score is 0.7747760448867355
```

Training the model

```
In [68]: clf2 = SGDClassifier(loss='hinge', n_jobs=-1, random_state=42, **r_search.best_params_)
clf2.fit(x_tr, y_train)
cal_clf2 = CalibratedClassifierCV(clf2, cv='prefit')
cal_clf2.fit(x_tr, y_train)
Out[68]: CalibratedClassifierCV(base_estimator=SGDClassifier(n_jobs=-1, random_state=42),
cv='prefit')
In [69]: y_pred_cv = cal_clf2.predict(x_cv)
y_prob_cv = cal_clf2.predict_proba(x_cv)
y_pred = cal_clf2.predict(x_te)
y_prob = cal_clf2.predict_proba(x_te)
```

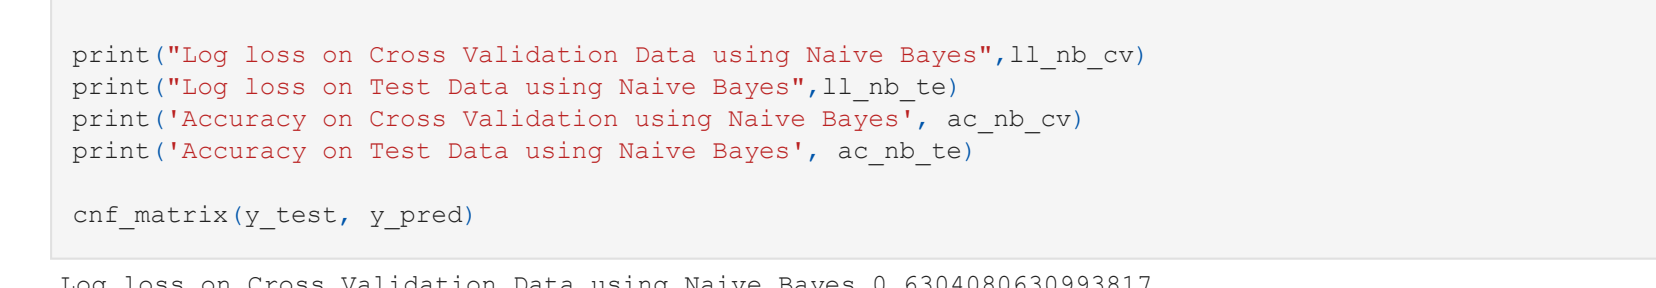
Performance of the model

```
In [70]: ll_svm_cv = log_loss(y_cv, y_prob_cv, eps=1e-15)
ac_svm_cv = accuracy_score(y_cv, y_pred_cv)
ll_svm_te = log_loss(y_test, y_prob, eps=1e-15)
ac_svm_te = accuracy_score(y_test, y_pred)
print('Log loss on Cross Validation Data using SVM', ll_svm_cv)
print('Log loss on Test Data using SVM', ll_svm_te)
print('Accuracy on Cross Validation using SVM', ac_svm_cv)
print('Accuracy on Test Data using SVM', ac_svm_te)
cmf_matrix(y_test, y_pred)
```


Log loss on Cross Validation Data using SVM 0.5183016042379512
Log loss on Test Data using SVM 0.53588640741728
Accuracy on Cross Validation using SVM 0.7429620563035496
Accuracy on Test Data using SVM 0.736532803696376



Precision matrix (Column Sum=1)




Recall matrix (Row sum=1)



Feature Importance

```
In [71]: importance = clf2.coef_
# summarize feature importance
for i, v in enumerate(importance[0]):
    print(f'Feature: {i}, Score: {v}')
plt.figure(figsize=(20,7))
sns.barplot(x=[x for x in range(importance.shape[1])], y=importance[0]).set_xticklabels(features, rotation=90)
plt.show()
Feature: 0, Score: 0.891598813947435
Feature: 1, Score: 0.9208021077456798
Feature: 2, Score: 0.0
Feature: 3, Score: 1.039811989140427
Feature: 4, Score: 0.80198932526858
Feature: 5, Score: -0.081958513947417
Feature: 6, Score: 0.0
Feature: 7, Score: 0.2970329379824765
Feature: 8, Score: -0.1485154689912395
Feature: 9, Score: 1.7524943499962597
Feature: 10, Score: 0.68175757359692
Feature: 11, Score: 1.22702576433421
Feature: 12, Score: 0.4755270077196574
Feature: 13, Score: 0.145861131754837
Feature: 14, Score: 0.234730028798932
Feature: 15, Score: 0.683175753597056
Feature: 16, Score: -1.34730028798932
Feature: 17, Score: 0.6534724635614665
Feature: 18, Score: 5.783585189417587
Feature: 19, Score: 0.2772670846800105
Feature: 20, Score: -0.03352627606561576
```



Naive Bayes

Hyperparameter Tuning

```
In [72]: alpha = [0.000001, 0.00001, 0.0001, 0.0001, 0.001, 0.01, 0.1, 1, 10, 50, 100]
params = {'alpha':alpha}
clf3 = MultinomialNB(alpha=alpha)
r_search = RandomizedSearchCV(clf3, param_distributions=params, return_train_score=True, random_state=42)
r_search.fit(x_tr, y_train)
Out[72]: RandomizedSearchCV(estimator=MultinomialNB(),
param_distributions={'alpha': [1e-07, 1e-06, 1e-05, 0.0001, 0.001, 0.01, 0.1, 1, 10, 50, 100]},
random_state=42, return_train_score=True)
In [73]: print(f'The best hyperparameter values is {r_search.best_params_} at which the score is {r_search.best_score}')
The best hyperparameter values is {'alpha': 0.01} at which the score is 0.670379797188602
```

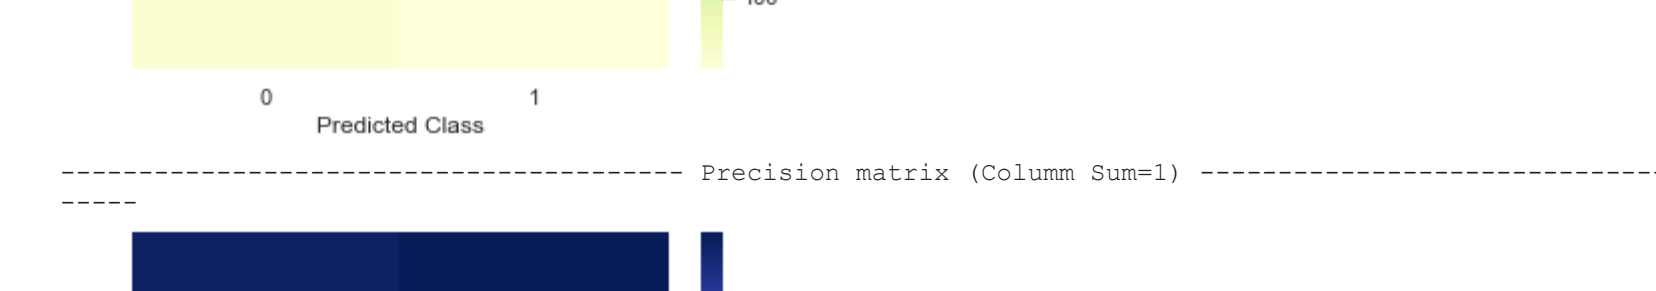
Training the model

```
In [74]: clf3 = MultinomialNB(**r_search.best_params_)
clf3.fit(x_tr, y_train)
cal_clf3 = CalibratedClassifierCV(clf3, cv='prefit')
cal_clf3.fit(x_tr, y_train)
Out[74]: CalibratedClassifierCV(base_estimator=MultinomialNB(alpha=0.01), cv='prefit')
In [75]: y_pred_cv = cal_clf3.predict(x_cv)
y_prob_cv = cal_clf3.predict_proba(x_cv)
y_pred = cal_clf3.predict(x_te)
y_prob = cal_clf3.predict_proba(x_te)
```

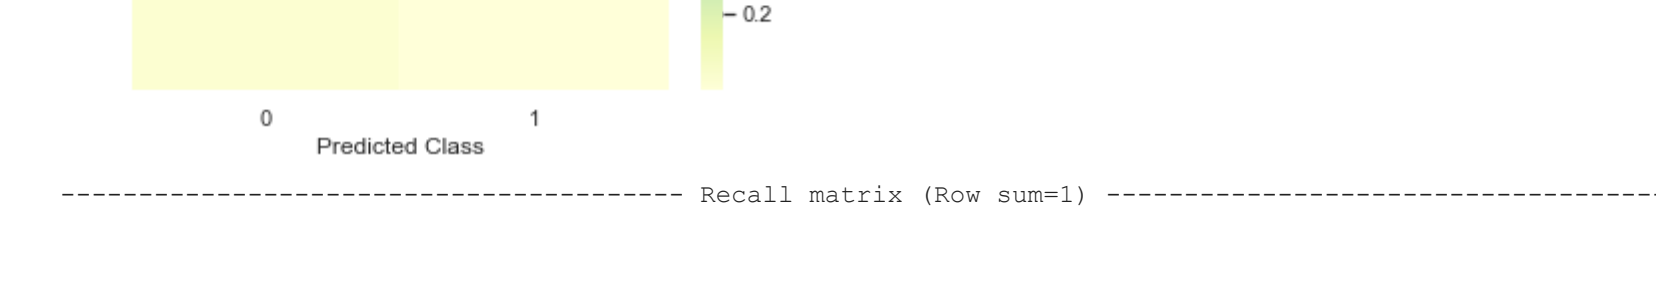
Performance of the model

```
In [76]: ll_nb_cv = log_loss(y_cv, y_prob_cv, eps=1e-15)
ac_nb_cv = accuracy_score(y_cv, y_pred_cv)
ll_nb_te = log_loss(y_test, y_prob, eps=1e-15)
ac_nb_te = accuracy_score(y_test, y_pred)
print('Log loss on Cross Validation Data using Naive Bayes', ll_nb_cv)
print('Log loss on Test Data using Naive Bayes', ll_nb_te)
print('Accuracy on Cross Validation using Naive Bayes', ac_nb_cv)
print('Accuracy on Test Data using Naive Bayes', ac_nb_te)
cmf_matrix(y_test, y_pred)
```

Log loss on Cross Validation Data using Naive Bayes 0.6304080630939817
Log loss on Test Data using Naive Bayes 0.399874013793987
Accuracy on Cross Validation using Naive Bayes 0.57723782129743
Accuracy on Test Data using Naive Bayes 0.616026263634868



Precision matrix (Column Sum=1)



Recall matrix (Row sum=1)

Random Forest using Sklearn

Hyperparameter Tuning

```
In [77]: # Maximum number of trees in the
n_estimators = [int(x) for x in np.linspace(start = 200, stop = 2000, num = 10)]
# Number of features to consider at every split
max_features = ['auto', None]
# Maximum number of nodes in the tree
max_depth = [int(x) for x in np.linspace(10, 110, num = 11)]
# Minimum number of samples required to split a node
min_samples_split = [2, 5, 10, 15]
# Minimum number of samples required at each leaf node
min_samples_leaf = [1, 2, 3, 4, 5]
# Method of selecting samples for training each tree
bootstrap = [True, False]
# Create the parameter space
params = {'n_estimators': n_estimators,
'max_features': max_features,
'max_depth': max_depth,
'min_samples_split': min_samples_split,
'min_samples_leaf': min_samples_leaf,
'bootstrap': bootstrap}
clf4 = RandomForestClassifier(n_estimators=100, max_depth=10, min_samples_split=10, min_samples_leaf=1, bootstrap=False)
r_search = RandomizedSearchCV(clf4, param_distributions=params, return_train_score=True, random_state=42)
r_search.fit(x_tr, y_train)
Out[77]: RandomizedSearchCV(estimator=RandomForestClassifier(n_estimators=100, max_depth=10, min_samples_split=10, min_samples_leaf=1, bootstrap=False),
param_distributions={'bootstrap': [True, False], 'max_depth': [10, 20, 30, 40, 50, 60, 70, 80, 90, 100, 110], 'max_features': ['auto', None], 'min_samples_split': [2, 5, 10, 15], 'min_samples_leaf': [1, 2, 3, 4, 5], 'n_estimators': [200, 400, 600, 800, 1000, 1200, 1400, 1600, 1800, 2000]},
random_state=42, return_train_score=True)
In [78]: print(f'The best hyperparameter values is {r_search.best_params_} at which the score is {r_search.best_score}')
The best hyperparameter values are {'n_estimators': 1400, 'min_samples_split': 10, 'min_samples_leaf': 1, 'max_features': 'auto', 'max_depth': 10, 'bootstrap': True} at which the score is 0.98968410223352
```

Training the model

```
In [79]: clf4 = RandomForestClassifier(**r_search.best_params_, n_jobs=-1, random_state=42)
clf4.fit(x_tr, y_train)
cal_clf4 = CalibratedClassifierCV(clf4, cv='prefit')
cal_clf4.fit(x_tr, y_train)
Out[79]: CalibratedClassifierCV(base_estimator=RandomForestClassifier(max_depth=10, max_features='log2', min_samples_split=10, n_estimators=1400, n_jobs=-1, random_state=42),
cv='prefit')
In [80]: y_pred_cv = cal_clf4.predict(x_cv)
y_prob_cv = cal_clf4.predict_proba(x_cv)
y_pred = cal_clf4.predict(x_te)
y_prob = cal_clf4.predict_proba(x_te)
```


Performance of the model

```
In [81]: ll_rf_cv = log_loss(y_cv, y_prob_cv, eps=1e-15)
ac_rf_cv = accuracy_score(y_cv, y_pred_cv)
ll_rf_te = log_loss(y_test, y_prob, eps=1e-15)
ac_rf_te = accuracy_score(y_test, y_pred)
print('Log loss on Cross Validation Data using RF', ll_rf_cv)
print('Log loss on Test Data using RF', ll_rf_te)
print('Accuracy on Cross Validation using RF', ac_rf_cv)
print('Accuracy on Test Data using RF', ac_rf_te)
cmf_matrix(y_test, y_pred)
```

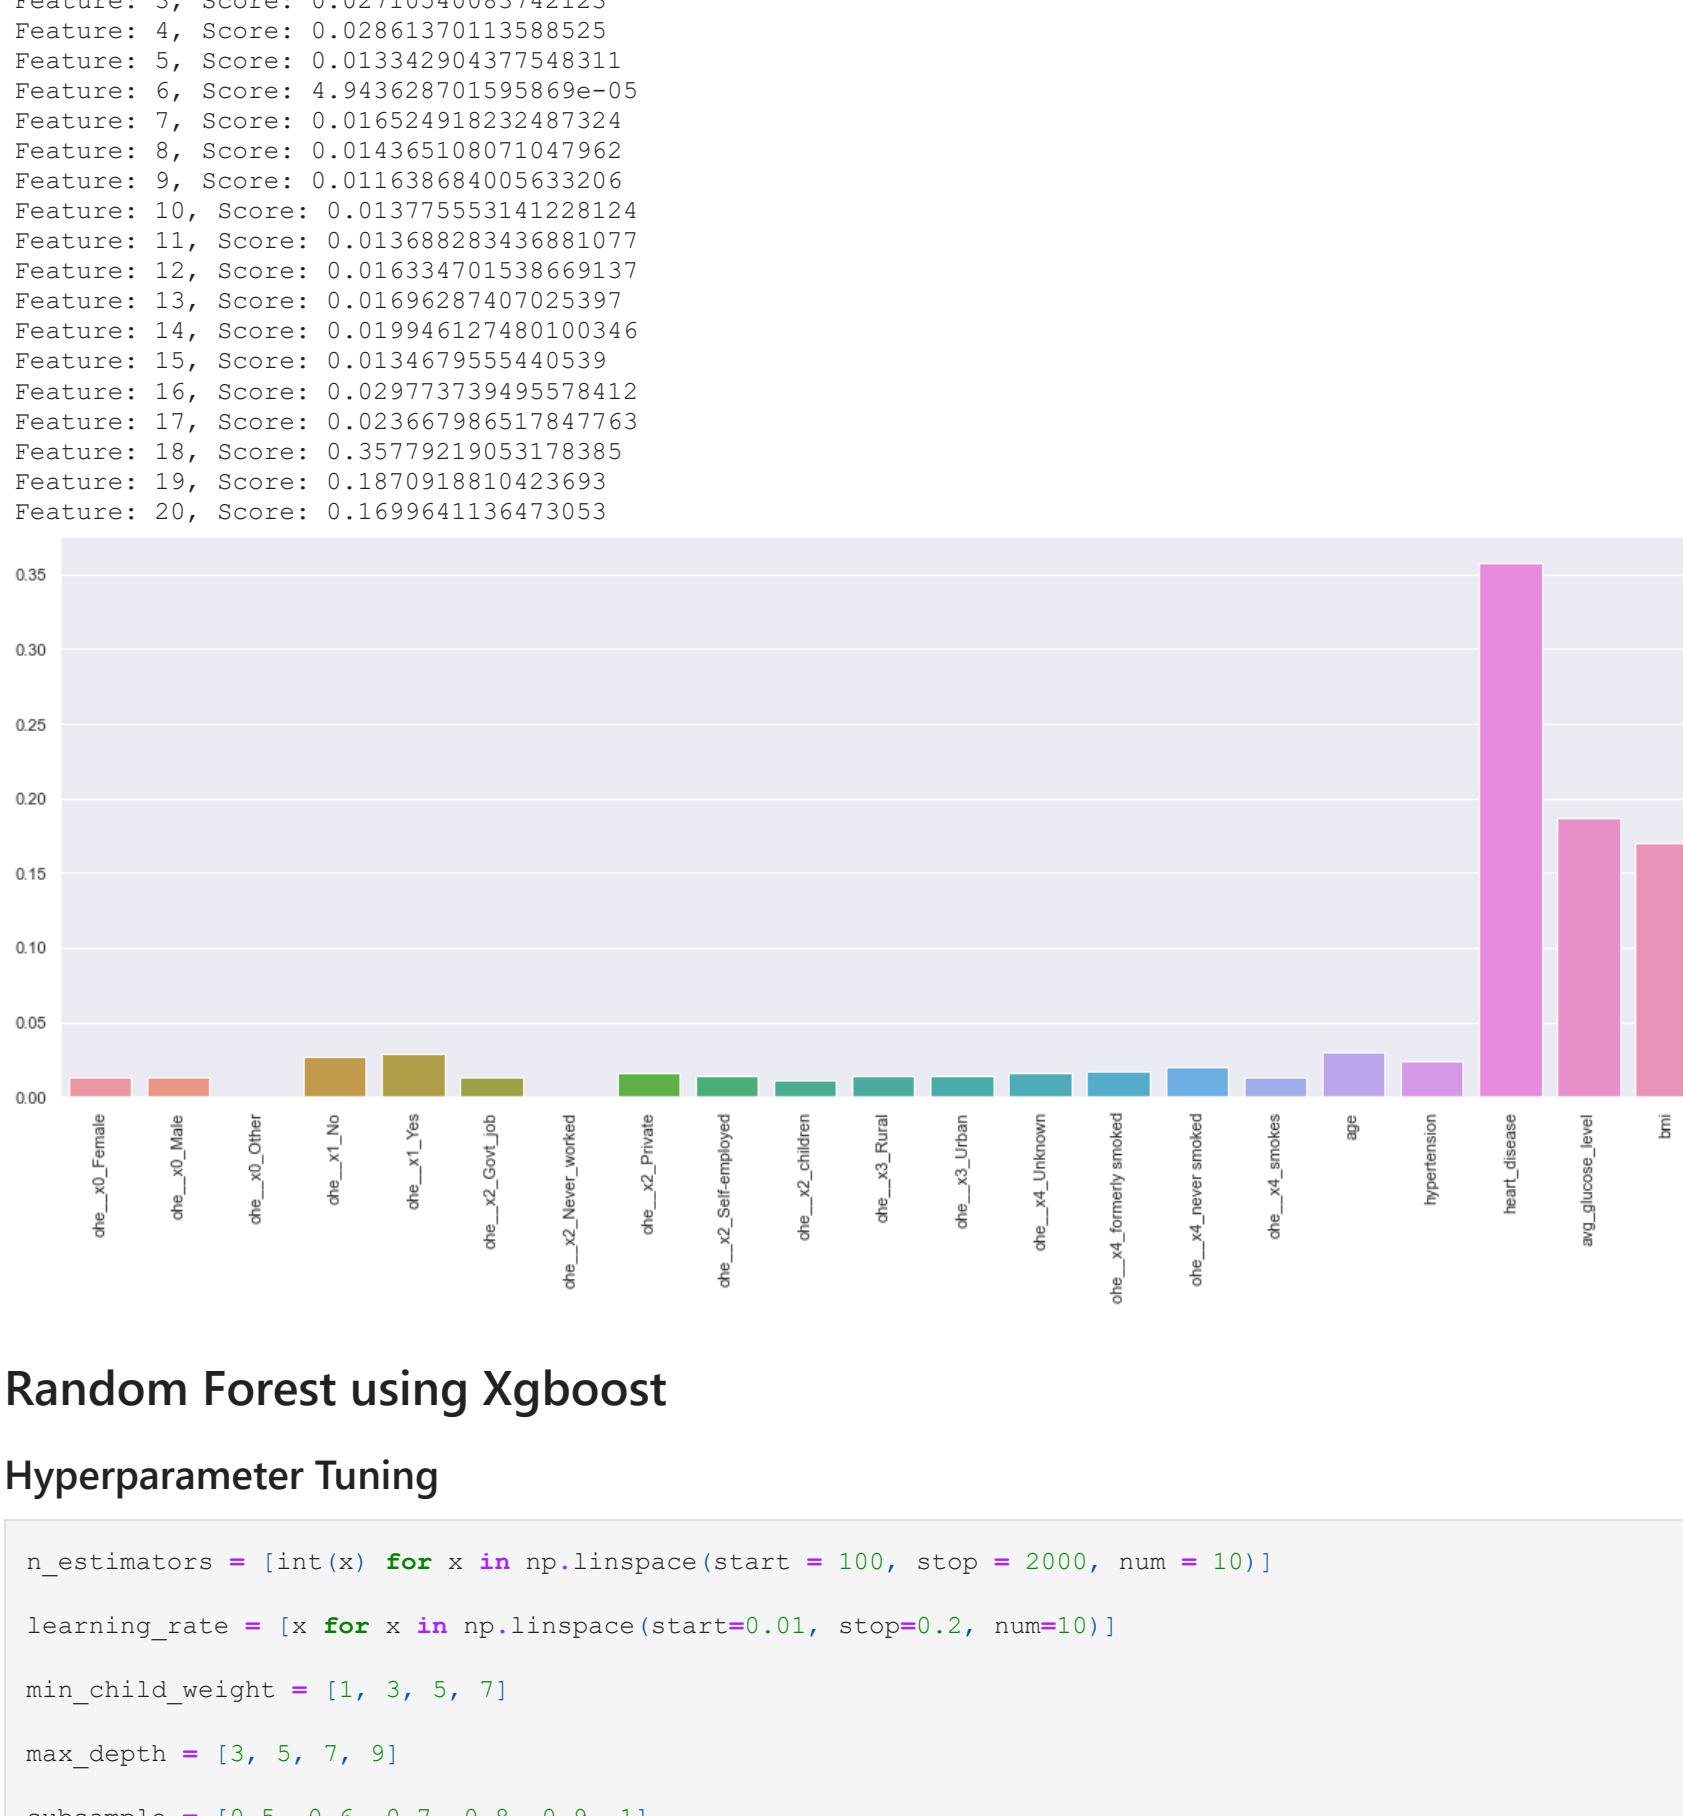
Log loss on Cross Validation Data using RF 0.3349050164010808
Log loss on Test Data using RF 0.3399874013793987
Accuracy on Cross Validation using RF 0.9424724602203183
Accuracy on Test Data using RF 0.94531813949095

Precision matrix (Column Sum=1)

Recall matrix (Row sum=1)

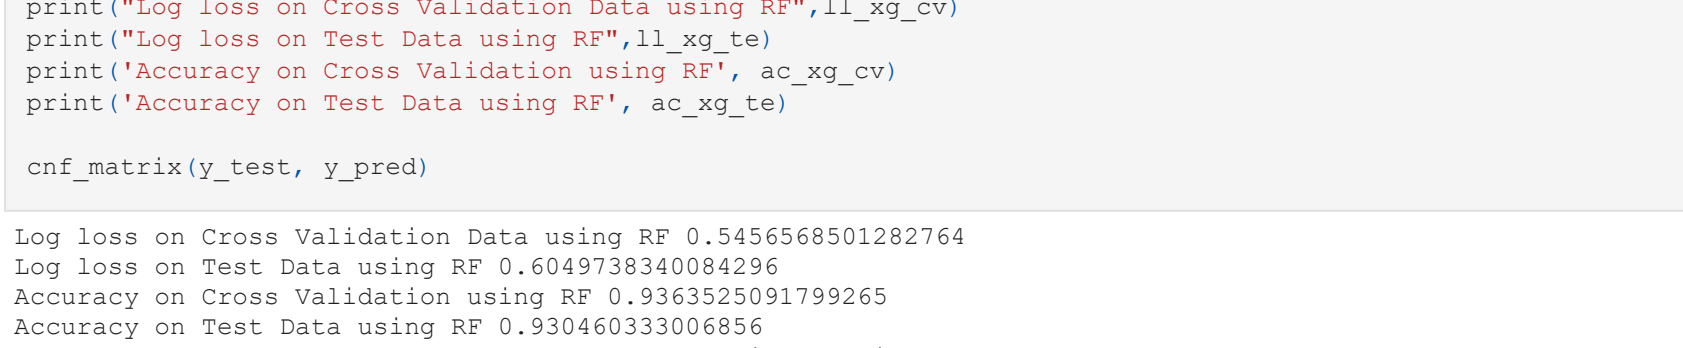



Feature Importance




Random Forest using Xgboost

Hyperparameter Tuning




The best hyperparameter values are {'subsample': 0.5, 'n_estimators': 100, 'min_child_weight': 1, 'max_depth': 5, 'learning_rate': 0.011111111111111111, 'colsample_bytree': 0.6} at which the score is 0.96283161458258

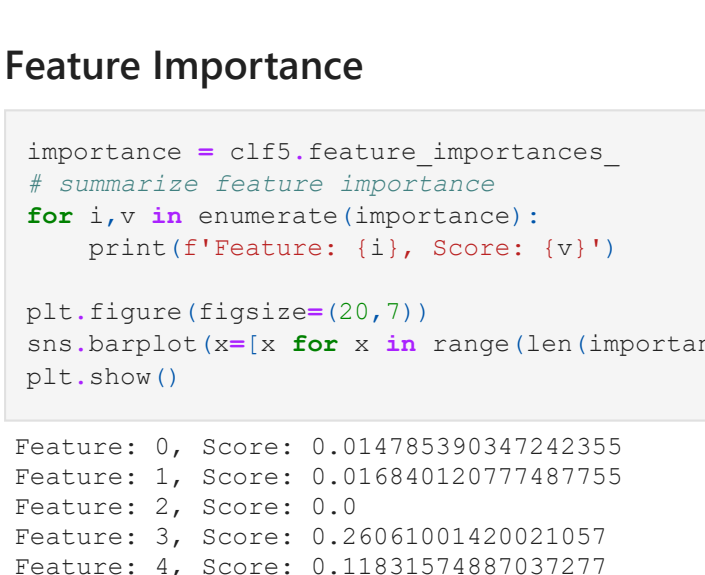
Training the model



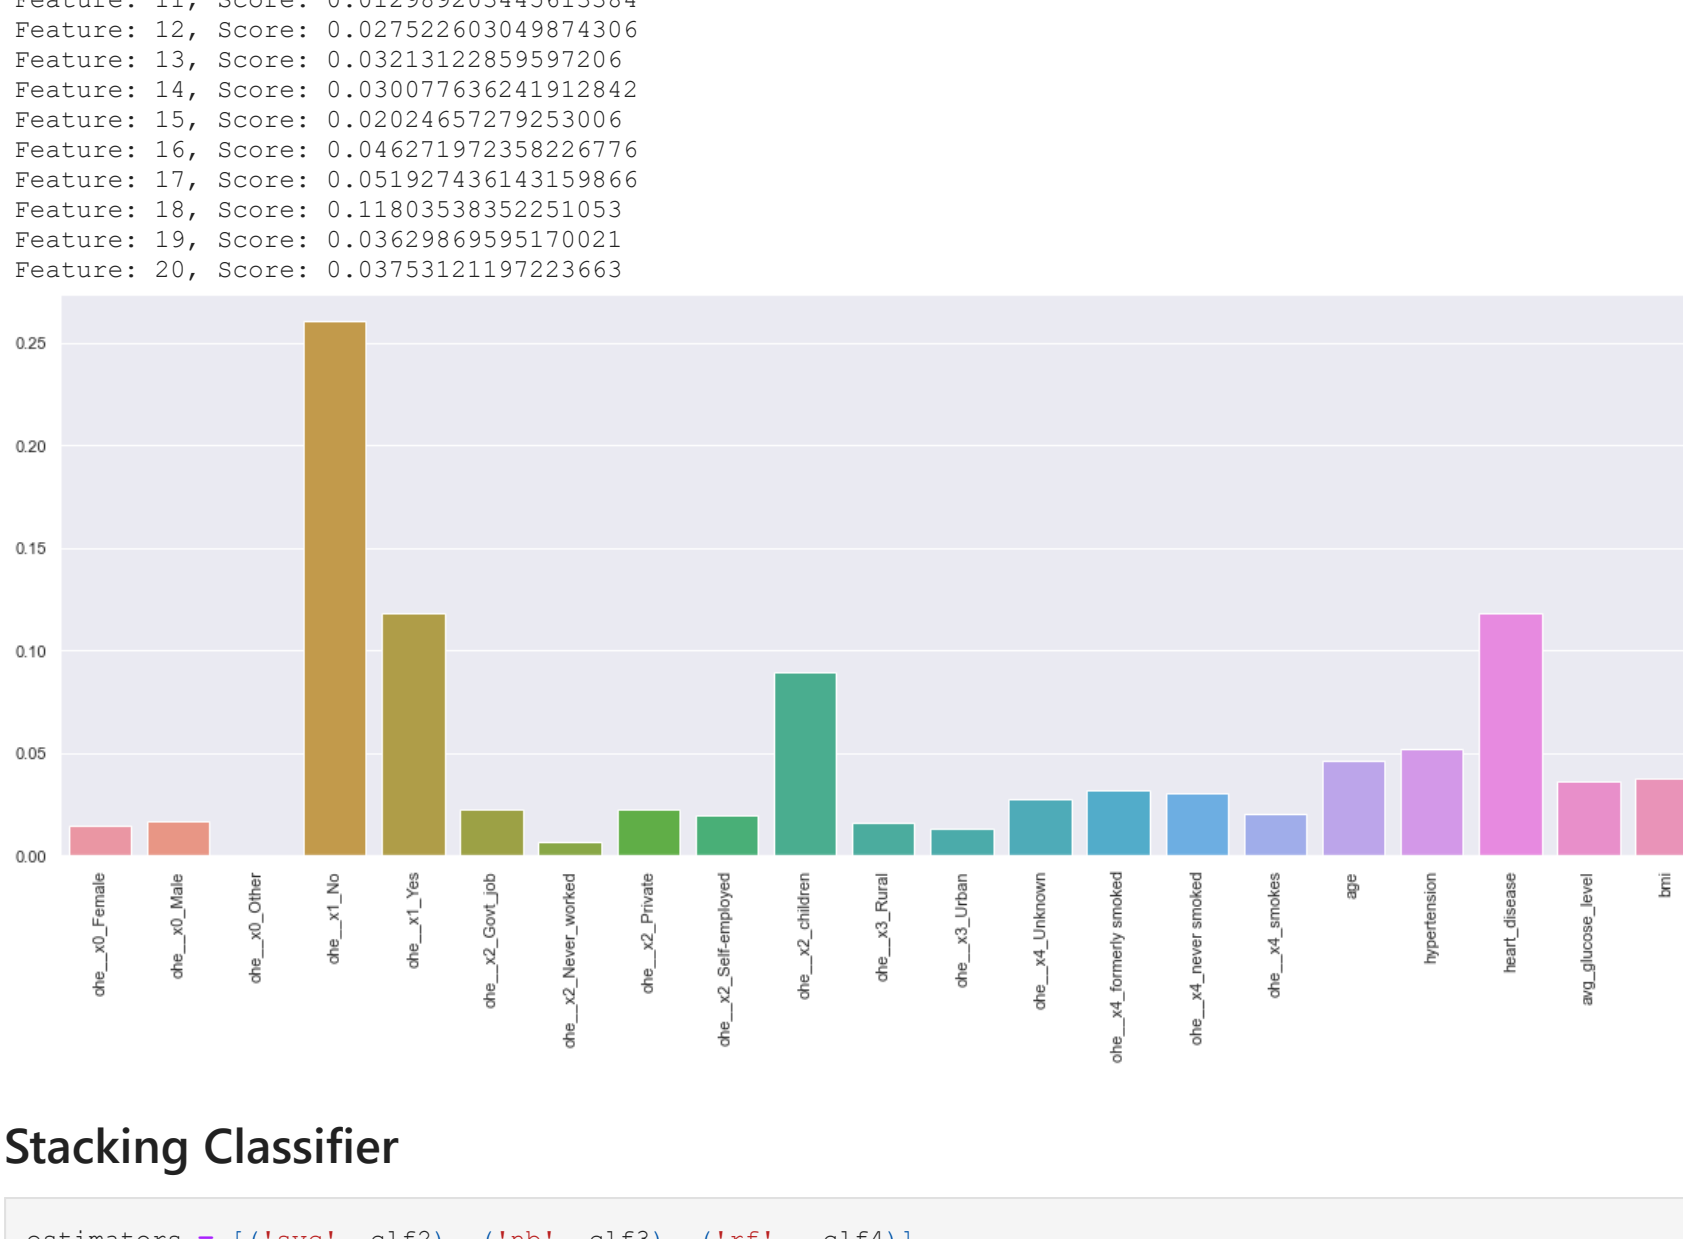
Performance of the model



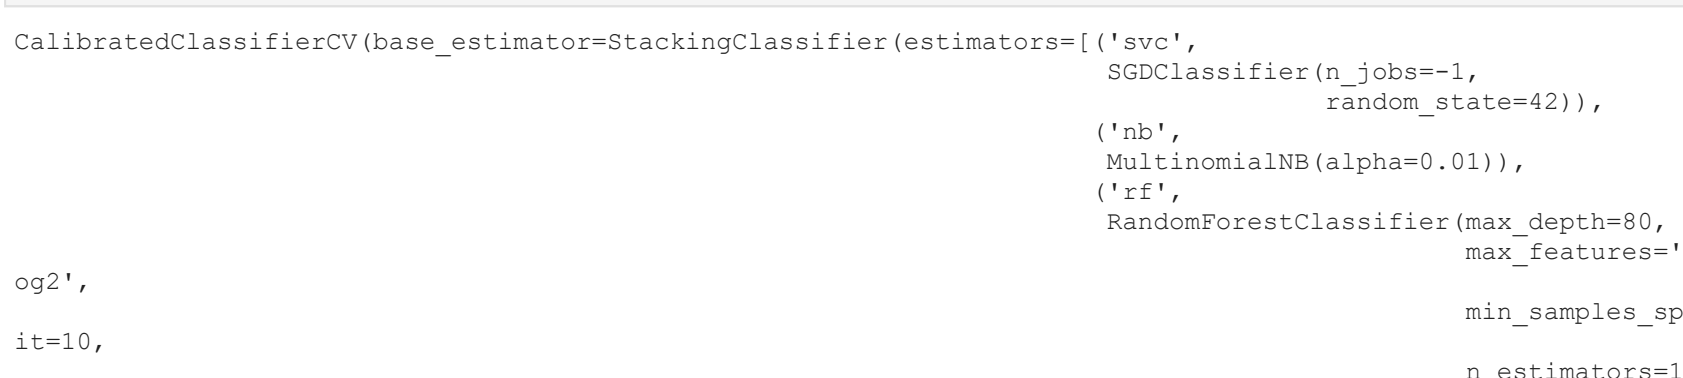
Log loss on Cross Validation Data using RF 0.545688801282764
Log loss on Test Data using RF 0.545688801282764
Accuracy on Cross Validation using RF 0.936352591793265
Accuracy on Test Data using RF 0.93646033006886



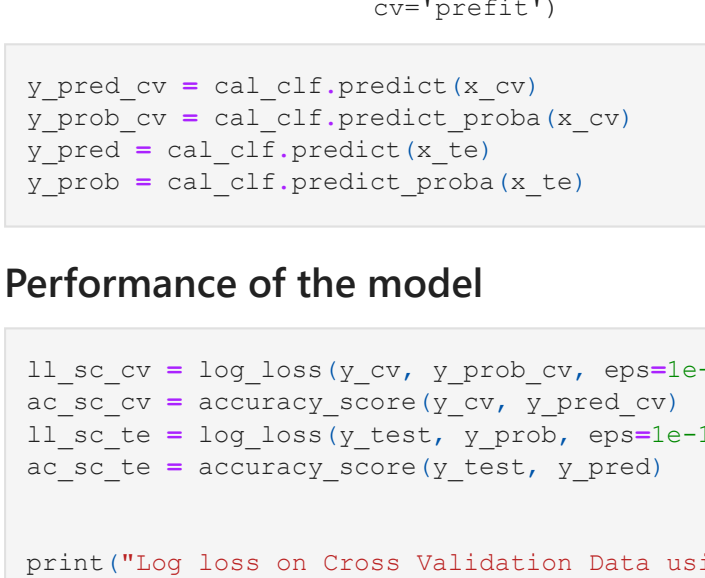
Feature Importance




Stacking Classifier




Log loss on Cross Validation Data using RF 0.545688801282764
Log loss on Test Data using RF 0.545688801282764
Accuracy on Cross Validation using RF 0.936352591793265
Accuracy on Test Data using RF 0.93646033006886




Performance of the model



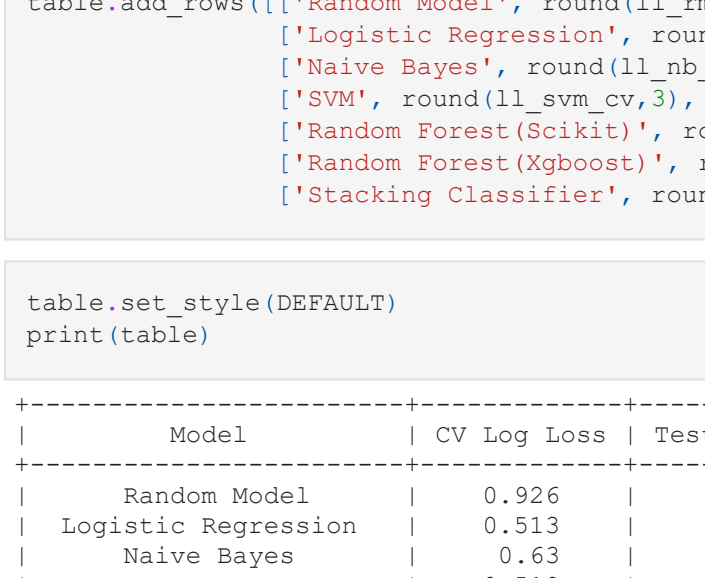
Log loss on Cross Validation Data using RF 0.545688801282764
Log loss on Test Data using RF 0.545688801282764
Accuracy on Cross Validation using RF 0.936352591793265
Accuracy on Test Data using RF 0.93646033006886



Summary of Performance



Log loss on Cross Validation Data using RF 0.545688801282764
Log loss on Test Data using RF 0.545688801282764
Accuracy on Cross Validation using RF 0.936352591793265
Accuracy on Test Data using RF 0.93646033006886



So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339

So our best model is Random Forest using sklearn with an accuracy of 94.5% and log loss of 0.339